

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

Commands + Code + Text

Files

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Shared drives

Disk 70.80 GB available

[1] from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

Necessary Imports
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
import matplotlib.pyplot as plt

[3] df = pd.read_csv("/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet2/Copy of mnist_dataset.csv")
print("Dataset Preview:")
print(df.head())
print("\nDataset Information:")
print(df.info())

Dataset Preview:
label pixel_0 pixel_1 pixel_2 pixel_3 pixel_4 pixel_5 pixel_6 \
0 5 0 0 0 0 0 0
1 0 0 0 0 0 0 0
2 4 0 0 0 0 0 0
3 1 0 0 0 0 0 0
4 9 0 0 0 0 0 0

pixel_7 pixel_8 ... pixel_774 pixel_775 pixel_776 pixel_777 \
0 0 0 ... 0 0 0 0
1 0 0 ... 0 0 0 0
2 0 0 ... 0 0 0 0
3 0 0 ... 0 0 0 0

114 358 879 390" data-label="Complex-Block">

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

Commands + Code + Text

Files

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...

Shared drives

Disk 70.80 GB available

> Softmax Function:
[] ↳ 1 cell hidden

> Softmax Test Case:
This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

[0] z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
softmax_output = softmax(z_test)
row_sums = np.sum(softmax_output, axis=1)
assert np.allclose(row_sums, 1), f"Test failed: Row sums are {row_sums}"
print("Softmax function passed the test case!")

Softmax function passed the test case!

114 555 879 587" data-label="Complex-Block">

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

Commands + Code + Text

Files

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...

Shared drives

Disk 70.80 GB available

> Test Function for Prediction Function:
The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

[0] X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

Expected Output:
The function should return an array with class labels (0, 1, or 2)

y_pred_test = predict_softmax(X_test, W_test, b_test)

Validate output shape
assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"

Print the predicted labels
print("Predicted class labels:", y_pred_test)

Predicted class labels: [1 1 0]

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

nmands + Code + Text

RAM Disk

les

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...

Loss Function:

[8] def loss_softmax(y_pred, y):
 """
 Compute the cross-entropy loss for a single sample.

 Parameters:
 y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,
 where c is the number of classes.
 y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of classes.

 Returns:
 float: Cross-entropy loss for the given sample.
 """
 # Avoid log(0) by adding a small epsilon (1e-12) to y_pred
 epsilon = 1e-12
 y_pred = np.clip(y_pred, epsilon, 1 - epsilon)

 loss = -np.sum(y * np.log(y_pred))

 return loss

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

nmands + Code + Text

RAM Disk

les

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions.

- Expects low loss for correct predictions.
- Expects high loss for incorrect predictions.

import numpy as np
y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]])
y_pred_correct = np.array([[0.9, 0.05, 0.05],
 [0.1, 0.85, 0.05],
 [0.05, 0.1, 0.85]])

Define incorrect predictions (high loss scenario)
y_pred_incorrect = np.array([[0.05, 0.05, 0.9],
 [0.1, 0.05, 0.85],
 [0.85, 0.1, 0.05]])

Compute loss for both cases
loss_correct = loss_softmax(y_pred_correct, y_true_correct)
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)

Validate that incorrect predictions lead to a higher loss
assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct < loss_incorrect, but got {loss_correct:.4f} >= {loss_incorrect:.4f}"

Print results
print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

nmands + Code + Text

RAM Disk

is

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Cost Function:

def cost_softmax(X, y, W, b):
 """
 Compute the average softmax regression cost (cross-entropy loss) over all samples.

 Parameters:
 X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d is the number of features.
 y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n is the number of samples and c is the number of classes.
 W (numpy.ndarray): Weight matrix of shape (d, c).
 b (numpy.ndarray): Bias vector of shape (c,).

 Returns:
 float: Average softmax cost (cross-entropy loss) over all samples.
 """
 n = X.shape[0] # Number of samples

 # Compute logits: XW + b
 logits = np.dot(X, W) + b # Shape: (n, c)

 # Apply softmax to get probabilities
 probabilities = softmax(logits) # Shape: (n, c)

 # Clip probabilities to avoid log(0)
 epsilon = 1e-12
 probabilities = np.clip(probabilities, epsilon, 1 - epsilon)

 # Compute cross-entropy loss for each sample
 losses = -np.sum(y * np.log(probabilities), axis=1) # Shape: (n,)

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

nands + Code + Text

Analyze your files with code written by Gemini Upload

RAM Disk

bs

[10]

Compute cross-entropy loss for each sample
losses = -np.sum(y * np.log(probabilities), axis=1) # Shape: (n,)

Compute total loss
total_loss = np.sum(losses)

Return average loss
return total_loss / n

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

nands + Code + Text

Analyze your files with code written by Gemini Upload

RAM Disk

bs

Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

import numpy as np

Example 1: Correct Prediction (Closer predictions)
X_correct = np.array([[1.0, 0.0], [0.0, 1.0]]) # Feature matrix for correct predictions
y_correct = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, matching predictions)
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]]) # Weights for correct prediction
b_correct = np.array([0.1, 0.1]) # Bias for correct prediction

Example 2: Incorrect Prediction (Far off predictions)
X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]]) # Feature matrix for incorrect predictions
y_incorrect = np.array([[1, 0], [0, 1]]) # True labels (one-hot encoded, incorrect predictions)
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]]) # Weights for incorrect prediction
b_incorrect = np.array([0.5, 0.6]) # Bias for incorrect prediction

Compute cost for correct predictions
cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)

Compute cost for incorrect predictions
cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect, b_incorrect)

Check if the cost for incorrect predictions is greater than for correct predictions
assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost {cost_incorrect} is not greater than correct cost {cost_correct}"

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

nands + Code + Text

Analyze your files with code written by Gemini Upload

RAM Disk

is

[12]

Print the costs for verification
print("Cost for correct prediction:", cost_correct)
print("Cost for incorrect prediction:", cost_incorrect)

print("Test passed!")

Cost for correct prediction: 0.0006234364133349324
Cost for incorrect prediction: 0.29930861359446115
Test passed!

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

inands + Code + Text

Analyze your files with code written by Gemini Upload

RAM Disk

les

Computing Gradients:

import numpy as np

def compute_gradient_softmax(X, y, W, b):
 """
 Compute the gradients of the cost function with respect to weights and biases.

 Parameters:
 X (numpy.ndarray): Feature matrix of shape (n, d).
 y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
 W (numpy.ndarray): Weight matrix of shape (d, c).
 b (numpy.ndarray): Bias vector of shape (c,).

 Returns:
 tuple: Gradients with respect to weights (d, c) and biases (c,).
 """

 # Compute logits: XW + b
 logits = np.dot(X, W) + b

 # Apply softmax to get probabilities
 y_pred = softmax(logits)

 # Compute the error term (difference between predicted and true labels)
 error = y_pred - y

 # Compute the gradient with respect to weights
 grad_W = np.dot(X.T, error) / X.shape[0] # Average over all samples (n, d)

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

nands + Code + Text

RAM Disk

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

Colab Notebooks

```
# Compute logits: XW + b
logits = np.dot(X, W) + b

# Apply softmax to get probabilities
y_pred = softmax(logits)

# Compute the error term (difference between predicted and true labels)
error = y_pred - y

# Compute the gradient with respect to weights
grad_W = np.dot(X.T, error) / X.shape[0] # Average over all samples (n, c)

# Compute the gradient with respect to biases
grad_b = np.sum(error, axis=0) / X.shape[0] # Average over all samples (c,)

return grad_W, grad_b
```

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

nands + Code + Text

RAM Disk

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

Colab Notebooks

- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- Untitled document.gd...
- week01-workshop1.ip...

Shared drives

70.80 GB available

Test case for compute_gradient function:

The test checks if the gradients from the function are close enough to the manually computed gradients using np.allclose, which accounts for potential floating-point discrepancies.

```
import numpy as np

# Define a simple feature matrix and true labels
X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]]) # Feature matrix (3 samples, 2 features)
y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]]) # True labels (one-hot encoded, 3 classes)

# Define weight matrix and bias vector
W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]]) # Weights (2 features, 3 classes)
b_test = np.array([0.1, 0.2, 0.3]) # Bias (3 classes)

# Compute the gradients using the function
grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)

# Manually compute the predicted probabilities (using softmax function)
z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

# Compute the manually computed gradients
grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

# Assert that the gradients computed by the function match the manually computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"
```

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

imands + Code + Text

RAM Disk

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

```
[14] # Assert that the gradients computed by the function match the manually computed gradients
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nExpected: {grad_W_manual}\nGot: {grad_W}"
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nExpected: {grad_b_manual}\nGot: {grad_b}"

# Print the gradients for verification
print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)

print("Test passed!")

Gradient w.r.t. W: [[ 0.1031051  0.01805685 -0.12116196]
 [-0.11608547  0.06079023  0.12921524]]
Gradient w.r.t. b: [-0.03290036  0.02484708  0.00805328]
Test passed!
```

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

les

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Shared Drives

Implementing Gradient Descent:

[15] def gradient_descent_softmax(X, y, W, b, alpha, n_iter, show_cost=False):
 """
 Perform gradient descent to optimize the weights and biases.

 Parameters:
 X (numpy.ndarray): Feature matrix of shape (n, d).
 y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).
 W (numpy.ndarray): Weight matrix of shape (d, c).
 b (numpy.ndarray): Bias vector of shape (c,).
 alpha (float): Learning rate.
 n_iter (int): Number of iterations.
 show_cost (bool): Whether to display the cost at intervals.

 Returns:
 tuple: Optimized weights, biases, and cost history.
 """
 cost_history = []

 for i in range(n_iter):
 # Compute gradients
 grad_W, grad_b = compute_gradient_softmax(X, y, W, b)

 # Update weights and biases using gradient descent
 W -= alpha * grad_W
 b -= alpha * grad_b

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Shared Drives

Update weights and biases using gradient descent
W -= alpha * grad_W
b -= alpha * grad_b

Compute the cost and add to cost history
cost = cost_softmax(X, y, W, b)
cost_history.append(cost)

Optionally display cost at intervals
if show_cost and i % 10 == 0:
 print(f"Iteration {i}, Cost: {cost:.4f}")

return W, b, cost_history

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

Worksheet2

Copy of 6CS012...
Copy of Iris.csv...
Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
2330545_RenjenSherpa...
Untitled document.gd...
week01-workshop1.ip...

Shared Drives

Preparing Dataset:

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def load_and_prepare_mnist(csv_file, test_size=0.2, random_state=42):
 """
 Reads the MNIST CSV file, splits data into train/test sets, and plots one image per class.

 Arguments:
 csv_file (str) : Path to the CSV file containing MNIST data.
 test_size (float) : Proportion of the data to use as the test set (default: 0.2).
 random_state (int) : Random seed for reproducibility (default: 42).

 Returns:
 X_train, X_test, y_train, y_test : Split dataset.
 """

 # Load dataset
 df = pd.read_csv(csv_file)

 # Separate labels and features
 y = df.iloc[:, 0].values # First column is the label
 X = df.iloc[:, 1:].values # Remaining columns are pixel values

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

Colab Notebooks

- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- Untitled document.gd...
- week01-workshop1.jp...

Shared drives

70.80 GB available

```
# Normalize pixel values (optional but recommended)
X = X / 255.0 # Scale values between 0 and 1

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=random_state)

# Plot one sample image per class
plot_sample_images(X, y)

return X_train, X_test, y_train, y_test

def plot_sample_images(X, y):
    """
    Plots one sample image for each digit class (0-9).

    Arguments:
    X (np.ndarray): Feature matrix containing pixel values.
    y (np.ndarray): Labels corresponding to images.
    """

    plt.figure(figsize=(10, 4))
    unique_classes = np.unique(y) # Get unique class labels

    for i, digit in enumerate(unique_classes):
        index = np.where(y == digit)[0][0] # Find first occurrence of the class
        image = X[index].reshape(28, 28) # Reshape 1D array to 28x28

        plt.subplot(2, 5, i + 1)
        plt.imshow(image, cmap='gray')
        plt.title(f"Digit: {digit}")
```

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

Colab Notebooks

- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- Untitled document.gd...
- week01-workshop1.jp...

Shared drives

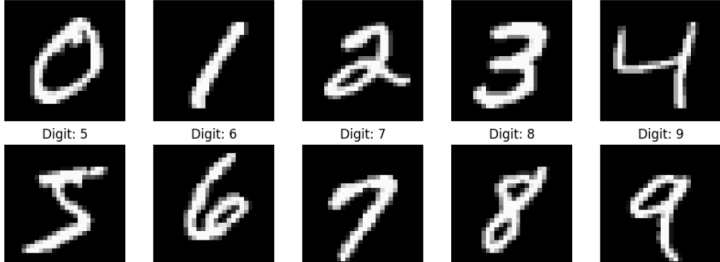
70.80 GB available

```
plt.imshow(image, cmap='gray')
plt.title(f"Digit: {digit}")
plt.axis('off')

plt.tight_layout()
plt.show()

csv_file_path = "/content/drive/myDrive/Artificial intelligence and machine learning/worksheet2/copy of mnist_dataset.csv" # Path to saved datas
X_train, X_test, y_train, y_test = load_and_prepare_mnist(csv_file_path)
```

Digit: 0 Digit: 1 Digit: 2 Digit: 3 Digit: 4



Digit: 5 Digit: 6 Digit: 7 Digit: 8 Digit: 9

RenjenSherpa_2330545_Worksheet02.json

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

Worksheet2

- Copy of 6CS012...
- Copy of Iris.csv...
- Copy of mnist_d...

Big Data

Colab Notebooks

- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- 2330545_RenjenSherpa...
- Untitled document.gd...
- week01-workshop1.jp...

Shared drives

70.80 GB available

A Quick debugging Step:

```
[19] # Assert that X and y have matching lengths
assert len(X_train) == len(y_train), f"Error: X and y have different lengths! X={len(X_train)}, y={len(y_train)}"
print("Move forward: Dimension of Feature Matrix X and label vector y matched.")
```

Move forward: Dimension of Feature Matrix X and label vector y matched.

Train the Model:

```
[20] print(f"Training data shape: {X_train.shape}")
print(f"Test data shape: {X_test.shape}")

Training data shape: (48000, 784)
Test data shape: (12000, 784)

[21] from sklearn.preprocessing import OneHotEncoder

# Check if y_train is one-hot encoded
if len(y_train.shape) == 1:
    encoder = OneHotEncoder(sparse_output=False) # Use sparse_output=False for newer versions of sklearn
    y_train = encoder.fit_transform(y_train.reshape(-1, 1)) # One-hot encode labels
    y_test = encoder.transform(y_test.reshape(-1, 1)) # One-hot encode test labels

# Now y_train is one-hot encoded, and we can proceed to use it
d = X_train.shape[1] # Number of features (columns in X_train)
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot encoding)
```

RenjenSherpa_2330545_Worksheet02.json ☆ 🌐

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

📁 ..

Worksheet2

Copy of 6CS012...

Copy of Iris.csv...

Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

Untitled document.gd...

week01-workshop1.ip...

Shared drives

70.80 GB available

[21] # Now y_train is one-hot encoded, and we can proceed to use it
d = X_train.shape[1] # Number of features (columns in X_train)
c = y_train.shape[1] # Number of classes (columns in y_train after one-hot encoding)

Initialize weights with small random values and biases with zeros
W = np.random.randn(d, c) * 0.01 # Small random weights initialized
b = np.zeros(c) # Bias initialized to 0

Set hyperparameters for gradient descent
alpha = 0.1 # Learning rate
n_iter = 1000 # Number of iterations to run gradient descent

Train the model using gradient descent
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b, alpha, n_iter, show_cost=True)

Plot the cost history to visualize the convergence
plt.plot(cost_history)
plt.title('Cost Function vs. Iterations')
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.grid(True)
plt.show()

Iteration 0, Cost: 2.2128
Iteration 10, Cost: 1.5064
Iteration 20, Cost: 1.1661
Iteration 30, Cost: 0.9817
Iteration 40, Cost: 0.8679
Iteration 50, Cost: 0.7906
Iteration 60, Cost: 0.7345
Iteration 70, Cost: 0.6917

RenjenSherpa_2330545_Worksheet02.json ☆ 🌐

File Edit View Insert Runtime Tools Help

mands + Code + Text

es

Analyze your files with code written by Gemini Upload

📁 ..

Worksheet2

Copy of 6CS012...

Copy of Iris.csv...

Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

Untitled document.gd...

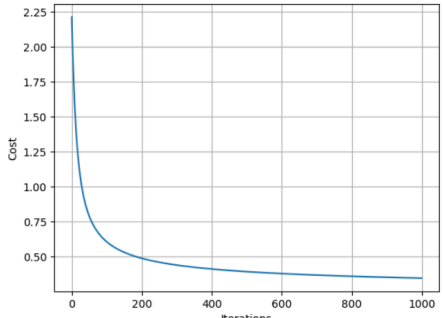
week01-workshop1.ip...

Shared drives

70.80 GB available

Iteration 810, Cost: 0.3531
Iteration 920, Cost: 0.3525
Iteration 930, Cost: 0.3519
Iteration 940, Cost: 0.3512
Iteration 950, Cost: 0.3506
Iteration 960, Cost: 0.3500
Iteration 970, Cost: 0.3494
Iteration 980, Cost: 0.3488
Iteration 990, Cost: 0.3483

Cost Function vs. Iterations



RenjenSherpa_2330545_Worksheet02.json ☆ 🌐

File Edit View Insert Runtime Tools Help

mands + Code + Text

les

Analyze your files with code written by Gemini Upload

📁 ..

Worksheet2

Copy of 6CS012...

Copy of Iris.csv...

Copy of mnist_d...

Big Data

Colab Notebooks

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

2330545_RenjenSherpa...

Untitled document.gd...

week01-workshop1.ip...

Shared drives

70.80 GB available

Evaluating the Model:

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score

Evaluation Function
def evaluate_classification(y_true, y_pred):
 """
 Evaluate classification performance using confusion matrix, precision, recall, and F1-score.

 Parameters:
 y_true (numpy.ndarray): True labels
 y_pred (numpy.ndarray): Predicted labels

 Returns:
 tuple: Confusion matrix, precision, recall, F1 score
 """
 # Compute confusion matrix
 cm = confusion_matrix(y_true, y_pred)

 # Compute precision, recall, and F1-score
 precision = precision_score(y_true, y_pred, average='weighted')
 recall = recall_score(y_true, y_pred, average='weighted')
 f1 = f1_score(y_true, y_pred, average='weighted')


 return cm, precision, recall, f1

mands + Code + Text

les

Analyze your files with
code written by Gemini

Upload



```
# Predict on the test set
y_pred_test = predict_softmax(X_test, W_opt, b_opt)

# Evaluate accuracy
y_test_labels = np.argmax(y_test, axis=1) # True labels in numeric form

# Evaluate the model
cm, precision, recall, f1 = evaluate_classification(y_test_labels, y_pred_test)

# Print the evaluation metrics
print("\nConfusion Matrix:")
print(cm)
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")

# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(12, 12))
cax = ax.imshow(cm, cmap='Blues') # Use a color map for better visualization

# Dynamic number of classes
num_classes = cm.shape[0]
ax.set_xticks(range(num_classes))
ax.set_yticks(range(num_classes))
ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])


# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
```

mands + Code + Text

es

Analyze your files with
code written by Gemini

Upload




```
# Add labels to each cell in the confusion matrix
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(cm) / 2 else 'black')

# Add grid lines and axis labels
ax.grid(False)
plt.title('Confusion Matrix', fontsize=14)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('Actual Label', fontsize=12)

# Adjust layout
plt.tight_layout()
plt.colorbar(cax)
plt.show()

Confusion Matrix:
[[127  0  5  2  3 11  9  2 13  3]
 [ 0 1274  7 11  1  5  1  4 18  1]
 [ 1 15 1029 15 20  4 26 24 33  7]
 [ 8  5 33 1048  1 56  8  8 31 21]
 [ 1  5  7  1 1092  0 11  4  4 51]
 [22 14 14 43 12 924 14  7 41 13]
 [ 7  2  9  1 10 16 1120  2 10  0]
 [ 7 27 24  5 15  3  0 1182  7 29]
 [ 9 27 14 34  9 31 13  6 1002 15]
 [ 8  6 10 18 43  9  0 30 10 1052]]

Precision: 0.90
Recall: 0.90
F1-Score: 0.90
```

RenjenSherpa_2330545_Worksheet02.json ☆

File Edit View Insert Runtime Tools Help

Hands + Code + Text

Linear Separability and Logistic Regression:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

def generate_data(linearly_separable=True, n_samples=200):
    X, y = make_classification(n_samples=n_samples, n_features=2, n_classes=2,
                              n_redundant=0, n_clusters_per_class=1, random_state=42)
    if not linearly_separable:
        X, y = make_classification(n_samples=n_samples, n_features=2, n_classes=2,
                                  n_redundant=0, n_clusters_per_class=2, random_state=42)
    return X, y

def plot_data(X, y, title="Data Distribution"):
    plt.figure(figsize=(8, 6))
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired, edgecolors="k", s=80)
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.title(title)
    plt.show()

def train_logistic_regression(X, y):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    model = LogisticRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy: {accuracy:.4f}")
    return model
```



Share


Gemini

imands + Code + Text

iles

Analyze your files with
code written by Gemini

Upload



return model

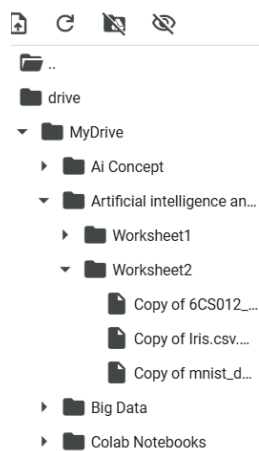
```
def plot_decision_boundary(X, y, model):  
    h = 0.02  
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1  
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1  
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))  
  
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])  
    Z = Z.reshape(xx.shape)  
  
    plt.figure(figsize=(8, 6))  
    plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.Paired)  
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors="k", cmap=plt.cm.Paired, s=80)  
    plt.xlabel("Feature 1")  
    plt.ylabel("Feature 2")  
    plt.title("Logistic Regression Decision Boundary")  
    plt.show()  
  
# Execute Steps  
X, y = generate_data(linearly_separable=True)  
plot_data(X, y, "Generated Data")  
  
model = train_logistic_regression(X, y)  
plot_decision_boundary(X, y, model)
```

imands + Code + Text


iles

Analyze your files with
code written by Gemini

Upload



```
model = train_logistic_regression(X, y)  
plot_decision_boundary(X, y, model)
```



Accuracy: 0.8750

les



Analyze your files with code written by Gemini

Upload



...

drive

MyDrive

AI Concept

Artificial intelligence an...

Worksheet1

Worksheet2

Copy of 6CS012_...

Copy of Iris.csv,...

Copy of mnist_d...

Big Data

Colab Notebooks

Logistic Regression Decision Boundary

