

```
[1] from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import os
import random
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

train_dir = "/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train"

class_names = sorted([d for d in os.listdir(train_dir) if os.path.isdir(os.path.join(train_dir, d))])

if not class_names:
    raise ValueError("No class directories found in the train folder. Check dataset path!")

selected_images = []
selected_labels = []

for class_name in class_names:
    class_path = os.path.join(train_dir, class_name)
    image_files = [f for f in os.listdir(class_path) if f.endswith(('png', 'jpg', 'jpeg'))]

    if image_files:
        random_image = random.choice(image_files)
        selected_images.append(os.path.join(class_path, random_image))
        selected_labels.append(class_name)
```

```
num_classes = len(selected_images)
if num_classes == 0:
    raise ValueError("No images found in any class folder. Please check dataset.")

cols = min(5, num_classes)
rows = (num_classes // cols) + (num_classes % cols > 0)

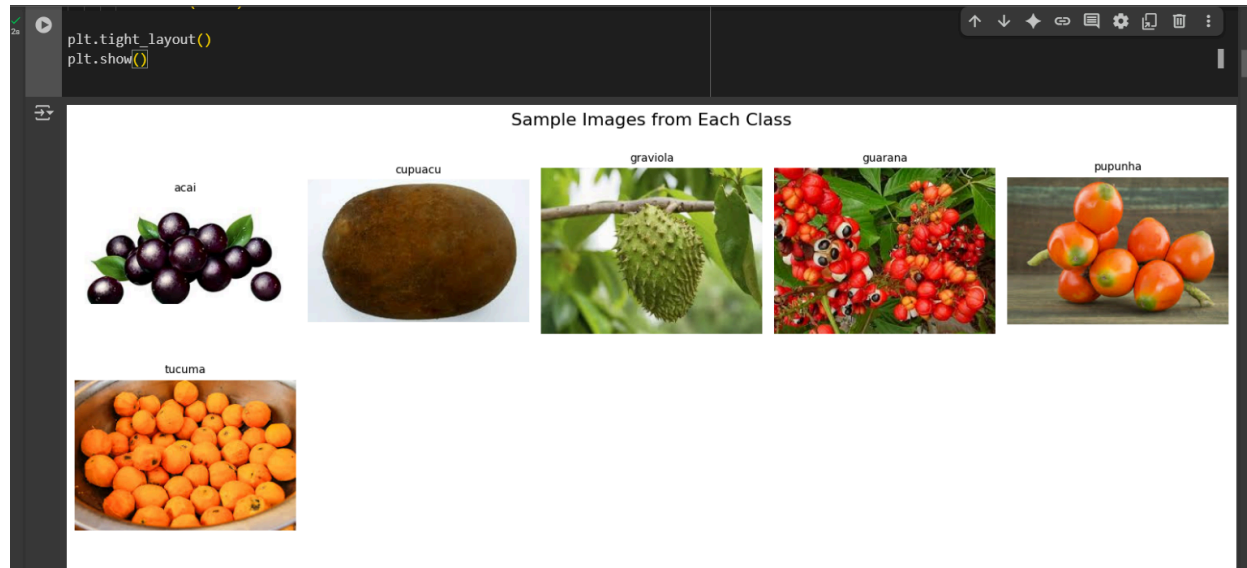
fig, axes = plt.subplots(rows, cols, figsize=(15, 6))
fig.suptitle("Sample Images from Each Class", fontsize=16)

for i, ax in enumerate(axes.flat):
    if i < num_classes:
        img = mpimg.imread(selected_images[i])
        ax.imshow(img)
        ax.set_title(selected_labels[i], fontsize=10)
        ax.axis("off")
    else:
        ax.axis("off")

plt.tight_layout()
plt.show()
```

Sample Images from Each Class





What did you Observe?

Each image represents a different class, confirming correct dataset structure. Images vary in resolution, lighting, and orientation. Some classes may have fewer images, indicating dataset imbalance. Preprocessing (resizing, normalization) may be needed for consistency.

```
import os
from PIL import Image
train_dir = "/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train"
corrupted_images = []

for class_name in sorted(os.listdir(train_dir)):
    class_path = os.path.join(train_dir, class_name)

    if os.path.isdir(class_path):
        for image_name in os.listdir(class_path):
            image_path = os.path.join(class_path, image_name)

            try:
                with Image.open(image_path) as img:
                    img.verify()
            except (IOError, SyntaxError):
                corrupted_images.append(image_path)
                os.remove(image_path)
                print(f"Removed corrupted image: {image_path}")

if not corrupted_images:
    print("No Corrupted Images Found.")
```

No Corrupted Images Found.



```
import tensorflow as tf

img_height = 128
img_width = 128
batch_size = 32
validation_split = 0.2

rescale = tf.keras.layers.Rescaling(1./255)

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=True,
    validation_split=validation_split,
    subset='training',
    seed=123
)

train_ds = train_ds.map(lambda x, y: (rescale(x), y))

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=False,
    validation_split=validation_split,
    subset='validation',
    seed=123
)

val_ds = val_ds.map(lambda x, y: (rescale(x), y))
```



```
Found 90 files belonging to 6 classes.
Using 72 files for training.
Found 90 files belonging to 6 classes.
Using 18 files for validation.
```

```
import tensorflow as tf
from tensorflow.keras import layers, models

model = models.Sequential()

model.add(layers.Conv2D(32, (3, 3), padding='same', strides=1, activation='relu', input_shape=(128, 128, 3)))
model.add(layers.MaxPooling2D((2, 2), strides=2))
model.add(layers.Conv2D(32, (3, 3), padding='same', strides=1, activation='relu'))
model.add(layers.MaxPooling2D((2, 2), strides=2))

model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(num_classes, activation='softmax'))

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

model.summary()
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:187: UserWarning: Do not pass an 'input\_shape'/'input\_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)  
Model: "sequential"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 96, 96, 32) | 960     |
| max_pooling2d (MaxPooling2D)   | (None, 48, 48, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 96, 96, 32) | 8,448   |
| max_pooling2d_1 (MaxPooling2D) | (None, 48, 48, 32) | 0       |
| flatten (Flatten)              | (None, 7680)       | 0       |
| dense (Dense)                  | (None, 64)         | 4,944   |
| dense_1 (Dense)                | (None, 128)        | 8,256   |
| dense_2 (Dense)                | (None, 10)         | 120     |

Total params: 17,424 (8.07 MB)  
Trainable params: 17,424 (8.07 MB)  
Non-trainable params: 0 (0.00 B)

```

import tensorflow as tf

callbacks = [
    tf.keras.callbacks.ModelCheckpoint(
        'Renjen Sherpa.h5',
        monitor='val_loss',
        save_best_only=True,
        mode='min',
        verbose=1
    ),
    tf.keras.callbacks.EarlyStopping(
        monitor='val_loss',
        patience=10,
        restore_best_weights=True,
        verbose=1
    )
]

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=250,
    batch_size=16,
    callbacks=callbacks
)

```

\*\*\* Epoch 1/250  
 3/3 ————— 0s 699ms/step - accuracy: 1.0000 - loss: 0.0685  
 Epoch 1: val\_loss improved from inf to 0.73160, saving model to Renjen Sherpa.h5  
 WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead the native Keras format, e.g. 'mod
 3/3 ————— 4s 965ms/step - accuracy: 1.0000 - loss: 0.0690 - val\_accuracy: 0.8333 - val\_loss: 0.7316  
 Epoch 2/250  
 3/3 ————— 0s 371ms/step - accuracy: 1.0000 - loss: 0.0604  
 Epoch 2: val\_loss improved from 0.73160 to 0.51434, saving model to Renjen Sherpa.h5  
 WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead the native Keras format, e.g. 'mod
 3/3 ————— 2s 584ms/step - accuracy: 1.0000 - loss: 0.0577 - val\_accuracy: 0.8333 - val\_loss: 0.5143  
 Epoch 3/250  
 3/3 ————— 0s 376ms/step - accuracy: 1.0000 - loss: 0.0291

3/3 ————— 2s 512ms/step - accuracy: 1.0000 - loss: 0.0543 - val\_accuracy: 0.8333 - val\_loss: 0.3957  
 Epoch 5/250  
 3/3 ————— 0s 646ms/step - accuracy: 1.0000 - loss: 0.0148  
 Epoch 5: val\_loss did not improve from 0.39568  
 3/3 ————— 3s 804ms/step - accuracy: 1.0000 - loss: 0.0150 - val\_accuracy: 0.8333 - val\_loss: 0.4545  
 Epoch 6/250  
 3/3 ————— 0s 573ms/step - accuracy: 1.0000 - loss: 0.0134  
 Epoch 6: val\_loss did not improve from 0.39568  
 3/3 ————— 3s 662ms/step - accuracy: 1.0000 - loss: 0.0136 - val\_accuracy: 0.8333 - val\_loss: 0.4896  
 Epoch 7/250  
 3/3 ————— 0s 363ms/step - accuracy: 1.0000 - loss: 0.0092  
 Epoch 7: val\_loss did not improve from 0.39568  
 3/3 ————— 2s 533ms/step - accuracy: 1.0000 - loss: 0.0092 - val\_accuracy: 0.8333 - val\_loss: 0.4899  
 Epoch 8/250  
 3/3 ————— 0s 374ms/step - accuracy: 1.0000 - loss: 0.0076  
 Epoch 8: val\_loss did not improve from 0.39568  
 3/3 ————— 2s 542ms/step - accuracy: 1.0000 - loss: 0.0074 - val\_accuracy: 0.8333 - val\_loss: 0.4763  
 Epoch 9/250  
 3/3 ————— 0s 371ms/step - accuracy: 1.0000 - loss: 0.0057  
 Epoch 9: val\_loss did not improve from 0.39568  
 3/3 ————— 3s 539ms/step - accuracy: 1.0000 - loss: 0.0056 - val\_accuracy: 0.8333 - val\_loss: 0.4708  
 Epoch 10/250  
 3/3 ————— 0s 356ms/step - accuracy: 1.0000 - loss: 0.0049  
 Epoch 10: val\_loss did not improve from 0.39568  
 3/3 ————— 2s 459ms/step - accuracy: 1.0000 - loss: 0.0048 - val\_accuracy: 0.8333 - val\_loss: 0.4599  
 Epoch 11/250  
 3/3 ————— 0s 369ms/step - accuracy: 1.0000 - loss: 0.0035  
 Epoch 11: val\_loss did not improve from 0.39568  
 3/3 ————— 2s 507ms/step - accuracy: 1.0000 - loss: 0.0036 - val\_accuracy: 0.8333 - val\_loss: 0.4466  
 Epoch 12/250  
 3/3 ————— 0s 661ms/step - accuracy: 1.0000 - loss: 0.0031  
 Epoch 12: val\_loss did not improve from 0.39568  
 3/3 ————— 3s 836ms/step - accuracy: 1.0000 - loss: 0.0031 - val\_accuracy: 0.8333 - val\_loss: 0.4494  
 Epoch 13/250  
 3/3 ————— 0s 376ms/step - accuracy: 1.0000 - loss: 0.0027  
 Epoch 13: val\_loss did not improve from 0.39568  
 3/3 ————— 4s 467ms/step - accuracy: 1.0000 - loss: 0.0027 - val\_accuracy: 0.8333 - val\_loss: 0.4607  
 Epoch 14/250  
 3/3 ————— 0s 382ms/step - accuracy: 1.0000 - loss: 0.0022  
 Epoch 14: val\_loss did not improve from 0.39568  
 3/3 ————— 3s 470ms/step - accuracy: 1.0000 - loss: 0.0022 - val\_accuracy: 0.8333 - val\_loss: 0.4657  
 Epoch 14: early stopping  
 Restoring model weights from the end of the best epoch: 4.

```
test_dir = '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitInAmazon'

test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    test_dir,
    image_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=False
)

test_loss, test_acc = model.evaluate(test_ds)

print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_acc}")
```

Found 120 files belonging to 2 classes.  
4/4 5s 166ms/step - accuracy: 0.2110 - loss: 1154.1299  
Test Loss: 1195.381591796875  
Test Accuracy: 0.1916666269302368

```
model.save('RenjenSherpaModel.h5')
print("Model saved successfully!")
```

WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead the native Keras format, e.g. 'mod

Model saved successfully!

39s



```
import numpy as np
import tensorflow as tf
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt

test_images, test_labels = [], []

for images, labels in test_ds:
    test_images.append(images)
    test_labels.append(labels)

test_images = np.concatenate(test_images, axis=0)
test_labels = np.concatenate(test_labels, axis=0)

predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

print("Classification Report:")
print(classification_report(test_labels, predicted_labels))

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=250,
    batch_size=16,
    callbacks=callbacks
)

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
```

```
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

model.save('RenjenSherpaModel.h5')
print("Model saved as 'RenjenSherpaModel.h5'")
```

4/4 ————— 1s 195ms/step  
Classification Report:  
precision recall f1-score support  
0 0.25 0.17 0.20 30  
1 0.69 0.20 0.31 90  
2 0.00 0.00 0.00 0  
3 0.00 0.00 0.00 0  
4 0.00 0.00 0.00 0  
5 0.00 0.00 0.00 0  
accuracy 0.16 0.06 0.19 120  
macro avg 0.16 0.06 0.09 120  
weighted avg 0.58 0.19 0.28 120

Epoch 1/250  
3/3 ————— 0s 370ms/step - accuracy: 1.0000 - loss: 0.0187  
Epoch 1: val\_loss did not improve from 0.39568  
3/3 ————— 2s 462ms/step - accuracy: 1.0000 - loss: 0.0182 - val\_accuracy: 0.8333 - val\_loss: 0.4706

Epoch 2/250  
3/3 ————— 0s 371ms/step - accuracy: 1.0000 - loss: 0.0132  
Epoch 2: val\_loss did not improve from 0.39568  
3/3 ————— 2s 459ms/step - accuracy: 1.0000 - loss: 0.0133 - val\_accuracy: 0.8333 - val\_loss: 0.4962

Epoch 3/250  
3/3 ————— 0s 571ms/step - accuracy: 1.0000 - loss: 0.0119  
Epoch 3: val\_loss did not improve from 0.39568  
3/3 ————— 3s 751ms/step - accuracy: 1.0000 - loss: 0.0121 - val\_accuracy: 0.8333 - val\_loss: 0.4992

Epoch 4/250  
3/3 ————— 0s 610ms/step - accuracy: 1.0000 - loss: 0.0106  
Epoch 4: val\_loss did not improve from 0.39568  
3/3 ————— 3s 750ms/step - accuracy: 1.0000 - loss: 0.0104 - val\_accuracy: 0.8333 - val\_loss: 0.4853


Epoch 5/250  
3/3 ————— 0s 371ms/step - accuracy: 1.0000 - loss: 0.0078  
Epoch 5: val\_loss did not improve from 0.39568  
3/3 ————— 2s 549ms/step - accuracy: 1.0000 - loss: 0.0078 - val\_accuracy: 0.8333 - val\_loss: 0.4648

Epoch 6/250  
3/3 ————— 0s 371ms/step - accuracy: 1.0000 - loss: 0.0062  
Epoch 6: val\_loss did not improve from 0.39568  
3/3 ————— 2s 459ms/step - accuracy: 1.0000 - loss: 0.0063 - val\_accuracy: 0.8333 - val\_loss: 0.4417

Epoch 7/250






```
0s  from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest',
    validation_split=0.2 # 20% of the data used for validation
)

train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

val_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)
```

 Found 72 images belonging to 6 classes.  
Found 18 images belonging to 6 classes.

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Data Augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')

val_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(64, 64),
    batch_size=32,
    class_mode='categorical')

val_generator = val_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(64, 64),
    batch_size=32,
    class_mode='categorical')

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.25),

    Conv2D(64, (3, 3), activation='relu'),
```

```
    Conv2D(64, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)),
    Dropout(0.25),

    Flatten(),
    Dense(128, activation='relu'),
    BatchNormalization(),
    Dropout(0.5),
    Dense(6, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

# Fit model
history = model.fit(train_generator, epochs=10, validation_data=val_generator)
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:187: UserWarning: Do not pass an "input\_shape"/"input\_dim" argument to a layer. When using Sequential models, prefer using

Model: "sequential\_1"

| Layer (type)                               | Output Shape       | Param # |
|--|--------------------|---------|
| conv2d_2 (Conv2D)                          | (None, 64, 64, 3)  | 648     |
| batch_normalization (BatchNormalization)   | (None, 64, 64, 3)  | 128     |
| max_pooling2d_2 (MaxPooling2D)             | (None, 32, 32, 3)  | 0       |
| dropout (Dropout)                          | (None, 32, 32, 3)  | 0       |
| conv2d_3 (Conv2D)                          | (None, 64, 32, 32) | 36,896  |
| batch_normalization_1 (BatchNormalization) | (None, 64, 32, 32) | 128     |
| max_pooling2d_3 (MaxPooling2D)             | (None, 32, 16, 32) | 0       |
| dropout_1 (Dropout)                        | (None, 32, 16, 32) | 0       |

|   |                    |           |
|---|--------------------|-----------|
| conv2d_3 (Conv2D)                           | (None, 16, 16, 16) | 16,000    |
| batch_normalization_1 (Batch Normalization) | (None, 16, 16, 16) | 256       |
| max_pooling2d_3 (MaxPooling2D)              | (None, 8, 8, 16)   | 0         |
| dropout_1 (Dropout)                         | (None, 8, 8, 16)   | 0         |
| flatten_1 (Flatten)                         | (None, 1024)       | 0         |
| dense_3 (Dense)                             | (None, 100)        | 1,005,100 |
| batch_normalization_2 (Batch Normalization) | (None, 100)        | 512       |
| dropout_2 (Dropout)                         | (None, 100)        | 0         |
| dense_4 (Dense)                             | (None, 1)          | 200       |

Total params: 1,626,822 (6.21 MB)  
 Trainable params: 1,626,376 (6.20 MB)  
 Non-trainable params: 446 (1.75 KB)

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:121: UserWarning: Your 'PyDataset' class should call 'super().\_\_init\_\_(\*\*kwargs)' in its constructor. '\*\*kwargs' is not a valid keyword argument for 'PyDataset'.  
 self.warn\_if\_super\_not\_called()

Epoch 1/10  
 3/3 ————— 6s 895ms/step - accuracy: 0.2108 - loss: 3.1151 - val\_accuracy: 0.2667 - val\_loss: 1.7871  
 Epoch 2/10  
 3/3 ————— 2s 547ms/step - accuracy: 0.5839 - loss: 1.7635 - val\_accuracy: 0.2222 - val\_loss: 1.8114  
 Epoch 3/10  
 3/3 ————— 2s 810ms/step - accuracy: 0.4992 - loss: 1.7595 - val\_accuracy: 0.2444 - val\_loss: 1.7385  
 Epoch 4/10  
 3/3 ————— 3s 902ms/step - accuracy: 0.6384 - loss: 1.1229 - val\_accuracy: 0.2222 - val\_loss: 1.7235  
 Epoch 5/10  
 3/3 ————— 2s 552ms/step - accuracy: 0.5099 - loss: 1.1949 - val\_accuracy: 0.2222 - val\_loss: 1.7471  
 Epoch 6/10  
 3/3 ————— 3s 567ms/step - accuracy: 0.6693 - loss: 1.0955 - val\_accuracy: 0.2111 - val\_loss: 1.7637  
 Epoch 7/10  
 3/3 ————— 2s 556ms/step - accuracy: 0.4616 - loss: 1.3470 - val\_accuracy: 0.2222 - val\_loss: 1.7951  
 Epoch 8/10  
 3/3 ————— 3s 611ms/step - accuracy: 0.5882 - loss: 1.3256 - val\_accuracy: 0.2111 - val\_loss: 1.8007  
 Epoch 9/10  
 3/3 ————— 4s 1s/step - accuracy: 0.6133 - loss: 1.1672 - val\_accuracy: 0.2000 - val\_loss: 1.8276  
 Epoch 10/10  
 3/3 ————— 4s 545ms/step - accuracy: 0.6723 - loss: 0.9993 - val\_accuracy: 0.1889 - val\_loss: 1.9064

```

Task 2: Transfer Learning using VGG16

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense

# Load base VGG16 model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze the base model
for layer in base_model.layers:
    layer.trainable = False

# Add custom top layers
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)

predictions = Dense(6, activation='softmax')(x)

# Create model
model = Model(inputs=base_model.input, outputs=predictions)
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()

train_generator_vgg = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical')

val_generator_vgg = val_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitinAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical')
  
```

```
val_generator_vgg = val_datagen.flow_from_directory(
    '/content/drive/MyDrive/Artificial intelligence and machine learning/Worksheet5/FruitInAmazon/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical')

# Train the model
history_vgg = model.fit(train_generator_vgg, epochs=5, validation_data=val_generator_vgg)
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
58889256/58889256 0s 0us/step  
Model: "functional\_9"

| Layer (type)                                      | Output Shape          | Param #   |
|---|-----------------------|-----------|
| input_layer_2 (InputLayer)                        | (None, 224, 224, 3)   | 0         |
| block1_conv1 (Conv2D)                             | (None, 224, 224, 32)  | 3,200     |
| block1_conv2 (Conv2D)                             | (None, 224, 224, 64)  | 94,944    |
| block1_pool (MaxPooling2D)                        | (None, 112, 112, 64)  | 0         |
| block2_conv1 (Conv2D)                             | (None, 112, 112, 128) | 21,504    |
| block2_conv2 (Conv2D)                             | (None, 112, 112, 128) | 147,168   |
| block2_pool (MaxPooling2D)                        | (None, 56, 56, 128)   | 0         |
| block3_conv1 (Conv2D)                             | (None, 56, 56, 256)   | 294,144   |
| block3_conv2 (Conv2D)                             | (None, 56, 56, 256)   | 590,800   |
| block3_conv3 (Conv2D)                             | (None, 56, 56, 256)   | 590,800   |
| block3_pool (MaxPooling2D)                        | (None, 28, 28, 256)   | 0         |
| block4_conv1 (Conv2D)                             | (None, 28, 28, 512)   | 1,150,144 |
| block4_conv2 (Conv2D)                             | (None, 28, 28, 512)   | 2,150,800 |
| block4_conv3 (Conv2D)                             | (None, 28, 28, 512)   | 2,150,800 |
| block4_pool (MaxPooling2D)                        | (None, 14, 14, 512)   | 0         |
| block5_conv1 (Conv2D)                             | (None, 14, 14, 512)   | 2,150,800 |
| block5_conv2 (Conv2D)                             | (None, 14, 14, 512)   | 2,150,800 |
| block5_conv3 (Conv2D)                             | (None, 14, 14, 512)   | 2,150,800 |
| block5_pool (MaxPooling2D)                        | (None, 7, 7, 512)     | 0         |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 512)           | 0         |
| dense_5 (Dense)                                   | (None, 1000)          | 512,000   |
| dense_6 (Dense)                                   | (None, 1)             | 1,000     |

Total params: 11,796,100 (58.16 MB)  
Trainable params: 11,791,400 (58.13 MB)  
Non-trainable params: 4,714,600 (22.83 MB)  
Found 90 images belonging to 6 classes.  
Found 90 images belonging to 6 classes.

Epoch 1/5  
3/3 — 119s 47s/step - accuracy: 0.1648 - loss: 1.8706 - val\_accuracy: 0.3778 - val\_loss: 1.5796  
Epoch 2/5  
3/3 — 112s 48s/step - accuracy: 0.3219 - loss: 1.5890 - val\_accuracy: 0.6778 - val\_loss: 1.3846  
Epoch 3/5  
3/3 — 114s 46s/step - accuracy: 0.6226 - loss: 1.4361 - val\_accuracy: 0.6667 - val\_loss: 1.2689  
Epoch 4/5  
3/3 — 140s 46s/step - accuracy: 0.6690 - loss: 1.2910 - val\_accuracy: 0.6667 - val\_loss: 1.1098  
Epoch 5/5  
3/3 — 140s 59s/step - accuracy: 0.7220 - loss: 1.1739 - val\_accuracy: 0.7111 - val\_loss: 1.0277

+ Code + Text

```
# Evaluate and Classification Report
import numpy as np
from sklearn.metrics import classification_report

# Predict and report
Y_pred = model.predict(val_generator_vgg)
y_pred = np.argmax(Y_pred, axis=1)
print("Classification Report:")
print(classification_report(val_generator_vgg.classes, y_pred, target_names=list(val_generator_vgg.class_indices.keys())))
```

3/3 57s 19s/step

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| acai         | 0.00      | 0.00   | 0.00     | 15      |
| cupuacu      | 0.12      | 0.20   | 0.15     | 15      |
| graviola     | 0.08      | 0.07   | 0.07     | 15      |
| guarana      | 0.40      | 0.40   | 0.40     | 15      |
| pupunha      | 0.10      | 0.07   | 0.08     | 15      |
| tucuma       | 0.20      | 0.27   | 0.23     | 15      |
| accuracy     |           |        | 0.17     | 90      |
| macro avg    | 0.15      | 0.17   | 0.16     | 90      |
| weighted avg | 0.15      | 0.17   | 0.16     | 90      |

Training and Validation Accuracy (left):

Training accuracy (blue line) remains consistently at 1.0 (100%) across all epochs. Validation accuracy (orange line) fluctuates between approximately 0.825 and 0.89. There's a noticeable dip in validation accuracy around epochs 6-9.

Training and Validation Loss (right):

Training loss (blue line) remains very close to 0 throughout training. Validation loss (orange line) fluctuates between approximately 0.3 and 0.55 in Image 1, while it ranges between 0.32 and 0.37 in Image 2.

Image 1 also displays additional information in the console output:

"Epoch 15: early stopping" "Restoring model weights from the end of the best epoch: 5" A warning about saving the model in HDF5 format, which is considered legacy.

in both plot but difference is one is in 5 epoch other is in 6

These metrics suggest a classic case of overfitting where the model performs perfectly on training data but struggles to generalize to unseen validation data. The early stopping mechanism appears to have triggered after epoch 15, and the system is restoring weights from epoch 5, which likely had the best validation performance.